

# tensormodels: Statistical Models for Tensor-Valued Data

by Anupama Kannan, David Kahle, and Michael Gallagher

**Abstract** *tensormodels* is an R package for simulation, maximum likelihood estimation, and model diagnostics for tensor-variate distributions arising in multiway array data. The package implements the tensor-variate normal distribution and several normal variance–mean mixture families, including the skew-t, generalized hyperbolic, variance gamma, and normal inverse Gaussian distributions. Estimation is performed via likelihood-based and expectation–maximization algorithms under a Kronecker-separable covariance model, yielding mode-specific covariance estimates through iterative updates. Model assessment tools include distance-based diagnostics, such as Mahalanobis distance plots constructed from tensor-whitened residuals. In addition, *tensormodels* provides implementations of widely used tensor decompositions for approximation and dimension reduction, including Tensor-Train decomposition via TT-SVD, Higher-Order Singular Value Decomposition (HOSVD), and Parallel Factor Analysis (CP/PARAFAC) computed using alternating least squares. It includes common linear algebra and tensor operations, such as matricization, the n-mode products, and the Khatri-Rao product.

## 1 Introduction

Lately, with the rise of data measurement tools and big data, the creation of higher dimensional data has become more and more common. These unique types of data require corresponding unique analytical tools to properly fit models with interpretation. Many researchers default to vectorizing or matricizing this higher order data to fit models that they are comfortable with and have pre-established tools. Additionally, they will commonly assume that data satisfies the assumptions of normality, which can sometimes be a very extreme choice that is only made to make computation and analysis easier.

The package **tensormodels** aims to assist with both of these issues. Firstly, it contains the tensor variate normal distribution and several skewed distributions under the normal mean-variance family model that can model a multitude of different shapes of data. Secondly, it contains diagnostic test and model heuristics to pick the best distribution for the user’s data set. It also contains random number generation and density functions for researchers who merely want to better understand tensor variate distributions.

In the following sections, we will explain the theory behind tensor variate distributions, compare related tensor software across several programming languages, explain some unique choices in the package design, show a simulation set up, and apply the models to a real-world problem.

## 2 Background and related work

The tensor space is still being established, with works using conflicting notation. In this section, we will set a standard for certain operations, and explain the basic underlying theory for several of the functions. Tensors are a generalization of vectors and matrices, allowing us to represent higher-order data into a structure. We need new terminology to describe these objects.

### 2.1 Order

The order of a tensor is the number of indices needed to identify an entry. This is the same as the number of values needed inside R’s subsetting brackets to pull one element out. A

vector is a first-order tensor, a matrix is a second-order tensor, and any object beyond that is a higher-order tensor.

For example, take the vector  $\mathbf{a} = (1, 2, 3)$ . To extract a specific element, you only require one value, so it has an order 1.

```
a <- c(1, 2, 3)
```

```
a[2]
#> [1] 2
```

The matrix

$$\mathbf{A} = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}.$$

has order two, since its entries are indexed by row and column.

```
A <- matrix(data = 1:6, nrow = 2, ncol = 3)
```

```
A[2, 1]
#> [1] 2
```

A third-order tensor requires three indices. The following tensor has dimensions  $4 \times 3 \times 2$ , so we need to specify a position along each mode to pull an element.




Figure 1: A  $4 \times 3 \times 2$  tensor  $\mathcal{A}$

```
A <- array(data = 1:24, dim = c(4, 3, 2))
```

```
A[3, 1, 1]
#> [1] 3
```

## 2.2 Mode

The mode is a specific axis of a tensor. For a matrix, the two modes are usually called rows and columns. However, since a tensor can have three or more dimensions, we now need a term for the specific dimension being referred to.

Consider a tensor  $\mathcal{A} \in 4 \times 2 \times 3$ . We would say the first mode has dimension 4, the second mode has dimension 2, and the third mode has dimension 3. Note that the modes of a tensor are dependent on the structure of the object. For example, we could keep the same data in  $\mathcal{A}$  but permute the modes, leading to a new tensor  $\mathcal{B} \in 4 \times 3 \times 2$ . When communicating a tensor, it is important to keep the modes consistent and keep track of what each mode represents. For example, the modes of a tensor could represent a specific sensor ID, time of measurement, and spatial location.

### 2.3 $n$ -mode product

In regular matrix multiplication, we can multiply two matrices as long as the number of columns in the first matrix matches the rows of the second matrix. Now, suppose we want to multiply a tensor and a matrix. Now, we need a new operator to specify which mode from the tensor should line up with the matrix. The  $n$ -mode product specifies that the  $n$ -th mode of a tensor should have the same size as the columns of a matrix.

Let tensor  $\mathcal{A} \in I_1 \times I_2 \times \dots \times I_n$  and let a matrix  $\mathbf{U} \in J \times I_{n+1} \times \dots \times I_N$ . Then the  $n$ -mode product of  $\mathcal{A}$  and  $\mathbf{U}$  can be defined as

$$(\mathcal{A} \times_n \mathbf{U}) \in I_1 \times \dots \times I_{n-1} \times J \times I_{n+1} \times \dots \times I_N.$$

Note that the modes of the result stay the same as  $\mathcal{A}$  other than the  $n$ -th mode being replaced with  $J$ .

Its entries are

$$(\mathcal{A} \times_n \mathbf{U})_{i_1, \dots, i_{n-1}, j, i_{n+1}, \dots, i_N} = \sum_{i_n=1}^{I_n} a_{i_1, i_2, \dots, i_N} u_{j, i_n}.$$

The tensor and matrix share one mode in common, denoted here as  $I_n$ . This operation is also called the tensor times matrix product.

This function is called `n_prod()` in the package **tenormodels**. It accepts a tensor, a matrix, and a value  $n$  specifying which mode to multiply across. For example, let us take  $\mathbf{a} \in \mathbb{R}^3$  and  $\mathbf{U} \in \mathbb{R}^{2 \times 3}$ . Then if we compute  $\mathbf{a} \times_1 \mathbf{U}$ , we expect the final result to be  $\in \mathbb{R}^{2 \times 1}$ .

```
a_vec <- c(1, 2, 3)
u_mat <- matrix(4:9, ncol = 3)

n_prod(tensor = a_vec, mat = u_mat, n = 1)
#> [1] 40 46
```

We can also do a more complex example with a tensor and a matrix. Here,  $\mathcal{A} \in \mathbb{R}^{2 \times 3 \times 4}$ . Then if we compute  $\mathcal{A} \times_2 \mathbf{U}$ , we expect the final result to be  $\in \mathbb{R}^{2 \times 2 \times 4}$ .

```
A_tensor <- array(1:24, dim = c(2, 3, 4))

dim(n_prod(A_tensor, u_mat, n = 2))
#> [1] 2 2 4
```

In **tenormodels**, the  $n$ -mode product is used internally to apply covariance transformations one mode at a time. For random generation, independent standard normal tensors are multiplied by Cholesky factors along each mode to induce the desired covariance structure. For density evaluation and model diagnostics, centered tensors are multiplied by inverse covariance factors along each mode, producing standardized tensor residuals and Mahalanobis distances.

### 2.4 Matricization

Matricization, also called unfolding, collapses a tensor into a matrix. It accepts an argument  $k$  that determines which mode to place in the rows, and then combines every other mode into the columns. This allows us to focus on variation along a specific mode.

Let  $\mathcal{A} \in I_1 \times \dots \times I_N$  be an  $N$ -th order tensor. The  $k$ -th matricization of  $\mathcal{A}$  is a matrix  $\mathbf{A}_k \in I_1 I_2 \dots I_k \times I_{k+1} \times I_{k+2} \times \dots \times I_N$ , whose elements are taken column-wise from  $\mathcal{A}$ . The elements are  $\mathbf{A}_{(k)}(i_1, \dots, i_k, i_{k+1}, \dots, i_N) = \mathcal{A}(i_1, \dots, i_k, i_{k+1}, \dots, i_N)$ .

In **tenormodels**, this function is called `matricization()`. It accepts a tensor and  $k$ .

### Example

The slices of  $A$  along the 3rd mode are two  $4 \times 3$  matrices.

$$A_1 = \begin{bmatrix} 1 & 5 & 9 \\ 2 & 6 & 10 \\ 3 & 7 & 11 \\ 4 & 8 & 12 \end{bmatrix}, \quad A_2 = \begin{bmatrix} 13 & 17 & 21 \\ 14 & 18 & 22 \\ 15 & 19 & 23 \\ 16 & 20 & 24 \end{bmatrix}.$$

Let us say we want to matricize along the first mode. This compresses the second and third mode into the columns of a matrix, so we would expect to obtain a matrix  $\in \mathbb{R}^{4 \times 6}$ . We will do this by moving the values in the third mode into the columns.

The notation for matricization along the first mode is  $A_{(1)}$ .

$$A_{(1)} = \begin{bmatrix} 1 & 5 & 9 & 13 & 17 & 21 \\ 2 & 6 & 10 & 14 & 18 & 22 \\ 3 & 7 & 11 & 15 & 19 & 23 \\ 4 & 8 & 12 & 16 & 20 & 24 \end{bmatrix}$$

```
A <- array(1:24, dim = c(4, 3, 2))
```

```
matricization(A, k = 1)
#>      [,1] [,2] [,3] [,4] [,5] [,6]
#> [1,]    1    5    9   13   17   21
#> [2,]    2    6   10   14   18   22
#> [3,]    3    7   11   15   19   23
#> [4,]    4    8   12   16   20   24
```

In **tenso**models, matricization is used to initialize and update mode-specific covariance matrices through ordinary matrix crossproducts.

## 2.5 Tensor variate distributions

There are several tensor variate distributions that are generalizations of matrix variate distributions. This section introduces the distributions that are used by **tenso**models.

The tensor variate normal distribution can be considered a generalization of the multivariate and matrix variate normal distributions.

### Univariate Normal

First, consider a standard univariate normal draw  $Z \sim N(0, 1)$ . If we shift and scale it to generate  $X = \mu + \sigma Z$ , then  $X \sim N(\mu, \sigma^2)$ . That is, an affine transformation of a standard univariate normal draw describes all possible univariate normal distributions.

### Multivariate Normal

That generated one element. Now let  $\mathbf{Z}$  be a  $p \times 1$  vector containing IID standard normal draws. Then  $\mathbf{Z} \sim N_p(\mathbf{0}, \mathbf{I}_p)$ . If we let  $\mathbf{A}\mathbf{A}^T = \Sigma$ , then define  $\mathbf{X} = \bar{\mathbf{x}} + \mathbf{A}\mathbf{Z}$ . Now  $\mathbf{X} \sim N_p(\bar{\mathbf{x}}, \Sigma)$ . That is,  $\mathbf{X}$  defines all the possible multivariate normal distributions.

### Matrix Variate Normal

Let  $\mathbf{Z}$  be a  $p \times q$  matrix with IID standard normal entries. Choose matrices  $\mathbf{A}$  and  $\mathbf{B}$  such that  $\mathbf{U} = \mathbf{A}\mathbf{A}^T$ ,  $\mathbf{V} = \mathbf{B}\mathbf{B}^T$ . Now define  $\mathbf{X} = \mathbf{M} + \mathbf{A}\mathbf{Z}\mathbf{B}^T$ . We say that  $\mathbf{X} \sim N_{pq}(\mathbf{M}, \mathbf{U}, \mathbf{V})$ . That is,  $\mathbf{X}$  defines all the possible matrix variate normal distributions.

**Table 1:** Mixing distributions used to construct the skewed tensor variate distributions.

| Distribution                           | Distribution of $W$                      | Parameters        | Parameter support                    |
|----------------------------------------|------------------------------------------|-------------------|--------------------------------------|
| Tensor variate skewed-t                | Inv-Gamma( $\nu/2, \nu/2$ )              | $\nu$             | $\nu > 0$                            |
| Tensor variate generalized hyperbolic  | Gen Inv Gaussian( $\omega, 1, \lambda$ ) | $\lambda, \omega$ | $\omega > 0, \lambda \in \mathbb{R}$ |
| Tensor variate variance gamma          | Gamma( $\gamma, \gamma$ )                | $\gamma$          | $\gamma > 0$                         |
| Tensor variate normal inverse Gaussian | Inv Gaussian( $1, \kappa$ )              | $\kappa$          | $\kappa > 0$                         |

### Tensor Variate Normal

Let  $\mathcal{Z}$  be an order- $D$  tensor with IID standard normal entries and  $\Sigma_1, \dots, \Sigma_D$  be positive definite matrices describing the variance-covariance structure along each mode. For each mode, choose  $\mathbf{A}_j$  such that  $\Sigma_j = \mathbf{A}_j \mathbf{A}_j^T$ ,  $j = 1, \dots, D$ .

Then we can define  $\mathcal{X}$  as  $\mathcal{X} = \mathcal{M} + \mathcal{Z} \times_1 \mathbf{A}_1 \times_2 \mathbf{A}_2 \cdots \times_D \mathbf{A}_D$ . That is,  $\mathcal{X}$  defines all the possible tensor variate normal distributions.

This is a nonlinear restriction on the natural parameter space of the full multivariate normal family. For the matrix variate normal, the covariance matrix is restricted to be defined as the Kronecker product of two covariance matrices. This makes its parameter space a lower-dimensional curved subset of the full multivariate normal distribution. (Roś et al., 2016)

### Skewed distributions as transformations of the tensor variate normal

To generate tensor variate skewed distributions, we use the normal variance mean mixture model (Barndorff-Nielsen et al., 1982). An  $r$ -dimensional random vector  $x$  is a normal variance-mean mixture with mixing distribution  $F$  if, for a given  $u \geq 0$  that follows a probability distribution  $F$  on  $[0, \infty)$ ,  $x|u \sim N_r(\mu + u\beta, u\Sigma)$ .  $\Sigma \in \mathbb{R}^{r \times r}$  is a constant, positive-definite matrix and  $\mu \in \mathbb{R}^r$  and  $\beta \in \mathbb{R}^r$  are constant vectors.

We state that  $x$  is a normal variance-mean mixture with position  $\mu$ , drift  $\beta$ , structure matrix  $\Sigma$ , and mixing distribution  $F$ . Different mixing distributions lead to different distributions. For example, if  $F$ , the mixing distribution, is the generalized inverse Gaussian distribution, then the resulting distribution of  $x$  is a generalized hyperbolic distribution.

In the tensor variate case, the normal variance mean mixture model is an efficient way to introduce skewness. A random tensor  $\mathcal{X}$  can be written as

$$\mathcal{X} = \mathcal{M} + W\mathcal{A} + \sqrt{W}\mathcal{V}, \quad (1)$$

where  $\mathcal{M}$  is a location tensor,  $\mathcal{A}$  is a skewness tensor,  $\mathcal{V} \sim \mathcal{N}_{n^*}(\mathbf{O}, \bigotimes_{d=1}^D \Sigma_d)$  and  $W > 0$  is a positive random variable. The distribution of  $\mathcal{X}$  will depend on the distribution of  $W$ .

The mixing distributions and their parameters for each of the skewed tensor variate distributions are summarized below. In each case, the support of  $W$  is  $(0, \infty)$ .

The tensor variate skewed-t distribution can be generated by setting  $W \sim \text{Inv-Gamma}(\nu/2, \nu/2)$  in Equation 1. We then say that  $\mathcal{X} \sim \text{TVST}_n(\mathcal{M}, \mathcal{A}, \bigotimes_{d=1}^D \Sigma_d, \nu)$ .

The tensor variate generalized hyperbolic distribution can be generated by setting  $W \sim I(\omega, 1, \lambda)$  in Equation 1. We then say that  $\mathcal{X} \sim \text{TVGH}_n(\mathcal{M}, \mathcal{A}, \bigotimes_{d=1}^D \Sigma_d, \lambda, \omega)$ .

The tensor variate variance gamma distribution can be generated by setting  $W \sim \text{Gamma}(\gamma, \gamma)$  in Equation 1. We then say that  $\mathcal{X} \sim \text{TVVG}_n(\mathcal{M}, \mathcal{A}, \bigotimes_{d=1}^D \Sigma_d, \gamma)$ .

The tensor variate normal inverse Gaussian distribution can be generated by setting  $W \sim \text{IG}(1, \kappa)$  in Equation 1. We then say that  $\mathcal{X} \sim \text{TVNIG}_n(\mathcal{M}, \mathcal{A}, \bigotimes_{d=1}^D \Sigma_d, \kappa)$ .

These choices of  $W$  give the four skewed families implemented in **tenso****models**. Their common representation also makes it possible to use a shared estimation framework while changing the distribution of the latent mixing variable.

### 3 Tensor software across programming languages

There are several existing software packages for tensors in several programming languages, but most emphasize tensor algebra or decompositions over probability models. The unique aspect of **tenso****models** is the support for tensor-variate normal and skewed tensor-variate distributions. These include random number generation, density functions, likelihood estimation, and model diagnostics.

One of the most popular softwares is the Tensor Toolbox for MATLAB (Bader et al., 2025; Bader and Kolda, 2006). This toolbox introduces a tensor class to generate dense tensors from a multidimensional MATLAB array. This class object has unique attributes so that the user can use methods such as `size`, `ndims`, `norm`, `unfold`, and `ttm`. It also supports decompositions on the tensor object with methods like `mttkrp`, `cp_als`, `tucker_als`, and `hosvd`. Tensor Toolbox also defines separate intentional classes for other tensor formats: `sptensor` for sparse tensors, `ktensor` for Kruskal/CP factored tensors, and `ttensor` for Tucker factored tensors (Bader and Kolda, 2007).

TensorLy in Python is similar to the Tensor Toolbox for MATLAB. It contains basic tensor operations and common decompositions including CP, Tucker, and Tensor-Train. A major proof of TensorLy is that it has backend support across other common frameworks in Python, including NumPy, PyTorch, and TensorFlow. This makes it easier for the user to run functions on their GPU or along multiple machines (Kossaifi et al., 2019). TensorFlow is another well-known package in Python, but it primarily uses the tensor object as inputs for machine learning methodologies rather than trying to expand on operations and decompositions.

The Python Tensor Toolbox, **pyttb** is a refactoring of the Tensor Toolbox for MATLAB. Thus, it ports the common functions from Tensor Toolbox into Python, including the various unique tensor classes and decompositions (pyttb developers, 2026).

In Julia, **TensorOperations.jl** have support for tensor contractions and **TensorToolbox.jl** is Julia's port of the MATLAB Tensor Toolbox (Periša, 2023; TensorOperations.jl developers, 2026).

In R, **rTensor** has an S4 Tensor class and includes common operations and decompositions. It contains the CP, Tucker, and HOSVD (Li et al., 2024a). **ttTensor** adds the Tensor-Train decomposition to this framework (Li et al., 2024b). **tenso****rr** is now archived on CRAN, but it originally provided support for S4 sparse tensor and dense tensor classes.

The comparison also explains the class design adopted in **tenso****models**. Packages such as MATLAB Tensor Toolbox and **rTensor** make the tensor class central: users create a tensor object first, and then methods dispatch on that object. **tenso****models** uses the base R array so that users do not have to learn the S4 class. This keeps **tenso****models** close to ordinary R array workflows.

| Purpose                              | tensormodels                                                                                                                                                         | Similar functions elsewhere                                                                           |
|--------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------|
| Create or store a dense tensor       | Uses base R array object                                                                                                                                             | MATLAB Tensor Toolbox 'tensor'; TensorLy 'tensor'; pyttb 'tensor'; rTensor 'Tensor'                   |
| Represent sparse or factored tensors | No separate sparse/factored tensor class; decomposition results are ordinary lists                                                                                   | MATLAB Tensor Toolbox 'sptensor', 'ktensor', 'ttensor'; pyttb equivalents; ITensors.jl 'ITensor'      |
| Matricize / unfold                   | 'matricization()'                                                                                                                                                    | MATLAB Tensor Toolbox 'unfold' / 'tenmat'; TensorLy 'unfold'; rTensor unfolding methods               |
| n-mode product                       | 'n_prod()', 'nm_prod()'                                                                                                                                              | MATLAB Tensor Toolbox 'ttm'; TensorLy 'mode_dot'; rTensor 'ttm'; TensorOperations.jl contractions     |
| CP/PARAFAC                           | 'cp()', 'cp_reconstruct()'                                                                                                                                           | MATLAB Tensor Toolbox 'cp_als'; Tensorlab 'cpd'; TensorLy 'parafac'; rTensor 'cp'                     |
| Tucker / HOSVD                       | 'tucker()', 'hosvd()', 'tucker_reconstruct()', 'hosvd_reconstruct()'                                                                                                 | MATLAB Tensor Toolbox 'tucker_als' and 'hosvd'; Tensorlab 'lmlra'; TensorLy 'tucker'; rTensor 'hosvd' |
| Tensor-Train                         | 'tt()', 'tt_reconstruct()'                                                                                                                                           | TensorLy 'tensor_train'; ttTensor 'TTSVD', 'TTCross', 'TTWOPT'                                        |
| Tensor-variate distributions         | 'rtnorm()', 'dtnorm()', 'rtskewt()', 'dtskewt()', 'rtvargamma()', 'dtvargamma()', 'rtgenhyper()', 'dtgenhyper()', 'rtinvgauss()', 'dtinvgauss()', and 'tensor_mle()' | Limited support in the general tensor software above                                                  |

## 4 Package design

The main goal of **tensormodels** is to provide functions for random number generation, density computation, and maximum likelihood estimation for common tensor variate distributions. Additionally, the package provides diagnostic tests for the structure of input data.

Distribution functions follow familiar R naming conventions: names beginning with *r* generate random draws, names beginning with *d* evaluate densities, and `tensor_mle()` fits a selected tensor-variate model. Tensor algebra helpers such as `matricization()`, `n_prod()`, `nm_prod()`, `tensor_prod()`, and `khatrao_rao()` are available to users and are also used internally by the likelihood and decomposition routines. Decomposition functions are grouped with explicit reconstruction helpers: `cp()` and `cp_reconstruct()`, `tucker()` and `tucker_reconstruct()`, `hosvd()` and `hosvd_reconstruct()`, and `tt()` and `tt_reconstruct()`.

One intentional decision was to store independent, identically distributed draws from a tensor variate distribution as a list. For example, if we generate  $n = 100$  draws from a tensor variate normal distribution of dimension  $2 \times 3 \times 4$ , the resulting R object will be a list of length 100. Each element in the list is a base R array object of dimension  $2 \times 3 \times 4$ . Other packages in this space created unique S3 and S4 class objects to store these types of draws. However, these choices are difficult for the average R user, as they have unique interactions and functions.

Another option would be to output a tensor of size  $100 \times 2 \times 3 \times 4$ ; that is, create a tensor of one higher dimension where the first mode represents the index of the draw. The choice to have the first mode represent the draw is arbitrary and could potentially confuse users, especially if they permute the dimensions.

```
three_dim_iid <- rtnorm(n = 100, mu = array(1:12, dim = c(2, 3, 4)))

length(three_dim_iid)
#> [1] 100
str(three_dim_iid[1:2], max.level = 1)
#> List of 2
#> $ : num [1:2, 1:3, 1:4] 1.52 3.22 2.57 3.85 6.34 ...
#> $ : num [1:2, 1:3, 1:4] -0.0797 1.5062 3.4586 3.817 5.5527 ...
#> - attr(*, "dim")= int 2
```

Model fits are returned as ordinary R lists. For example, a tensor-variate normal fit contains the estimated mean array  $\mu$ , the mode-specific covariance matrices  $\sigma$ s, likelihood-based summaries, and convergence information when relevant. Skewed families add distribution-specific parameters and skewness arrays. This return design keeps the model output easy to inspect with base R tools while avoiding a separate hierarchy of model classes.

#### 4.1 Core workflow

This package contains random number generations for several tensor variate distributions. Below is a simulation from a tensor variate normal of size  $2 \times 3 \times 4$ . The dimension of the draws is the same as the  $\mu$  object provided. The covariance structure is specified with a list of matrices, where each matrix must be a square of the corresponding mode.

```
mu_true <- array(1:24, dim = c(2, 3, 4))

S1 <- crossprod(matrix(rnorm(4), nrow = 2))
S2 <- crossprod(matrix(rnorm(9), nrow = 3))
S3 <- crossprod(matrix(rnorm(16), nrow = 4))

sigmas_true <- list(S1, S2, S3)

tvn_draws <- rtnorm(n = 1e3, mu = mu_true,
                   sigmas = sigmas_true)
```

Given a list of draws, the package can estimate the MLEs using the function `tensor_mle()`. The model argument is required to know which distribution to fit to the data.

```
tvn_mles <- tensor_mle(tvn_draws, model = "normal")

str(tvn_mles)
#> List of 6
#> $ mu      : num [1:2, 1:3, 1:4] 0.838 2.514 2.803 4.669 4.983 ...
#> $ sigmas:List of 3
#> ..$ : num [1:2, 1:2] 0.277 -0.383 -0.383 1.723
#> ..$ : num [1:3, 1:3] 1.389 1.416 0.119 1.416 1.597 ...
#> ..$ : num [1:4, 1:4] 17.53 3.23 3.36 4.64 3.23 ...
#> $ loglik: num -8984
#> $ k      : num 41
#> $ AIC    : num 18050
#> $ BIC    : num 18252
```

There are also density functions for each distribution. For the tensor variate normal, the required arguments are  $\mu$  and  $\sigma$ s. There is an optional `log` argument to compute the log density rather than the density itself.

```
dtnorm(tvn_draws[[1]], mu = mu_true, sigmas = sigmas_true)
#> [1] 0.001097981
```

#### 4.2 Tensor decompositions

The idea of decompositions can also apply to tensors. The package provides the Tucker, PARAFAC, higher-order singular value, and Tensor-Train decompositions.



### Tucker decomposition

The idea behind the Tucker decomposition is similar to rank factorization for matrices. The Tucker decomposition factorizes an  $N$ -th order tensor into a core tensor and a list of matrices ([Hitchcock, 1927](#)):

$$\mathcal{A} = \mathcal{B} \times_1 \mathbf{U}_1 \times_2 \mathbf{U}_2 \cdots \times_N \mathbf{U}_N.$$

The core tensor  $\mathcal{B}$  has the same order as the original tensor  $\mathcal{A}$ , and the number of  $\mathbf{U}$  matrices is also the number of modes. The modes of  $\mathcal{B}$  are called the Tucker ranks, which determine how reduced the form is.

Each matrix  $\mathbf{U}_n$  contains principal factors for mode  $n$ . The  $\mathbf{U}_n$  is a basis for the  $n$ th mode, and thus spans a subspace of  $\mathbb{R}^{I_n}$ . The core tensor  $\mathcal{B}$  represents the strength of interactions between factors. The idea behind this decomposition is to reduce the dimensionality of the input. We are able to summarize the higher-order object  $\mathcal{A}$  with a smaller core tensor  $\mathcal{B}$  and several factor matrices  $\mathbf{U}_1, \mathbf{U}_2, \dots, \mathbf{U}_N$ . However, one key issue with the Tucker decomposition is that it is not unique.

### PARAFAC decomposition

The PARAFAC decomposition has several other names—CANDECOMP, CP, and Tensor Rank decomposition. PARAFAC stands for parallel factor analysis while CANDECOMP stands for canonical polyadic decomposition. This decomposition is a special case of the Tucker decomposition, where the core tensor is superdiagonal.

Let us say we want to approximate the tensor  $\mathcal{A} \in I_1 \times I_2 \times \cdots \times I_N$  with a linear combination of  $R$  rank 1 tensors. The approximation can be written as

$$\mathcal{A} \approx \sum_{r=1}^R \lambda(r) \mathbf{u}_1(:, r) \otimes \cdots \otimes \mathbf{u}_N(:, r).$$

However, determining the rank of a tensor is an NP hard problem, so it is difficult to determine the value of  $R$  necessary to approximate the original tensor. The package computes this decomposition using the alternating least squares algorithm described by [Carroll and Chang \(1970\)](#) and [Harshman \(1970\)](#).

### Higher-order singular value decomposition

The higher-order singular value decomposition (HOSVD) is a specific case of the Tucker decomposition with restrictions. The matrices  $\mathbf{U}$  are unitary and the core tensor  $\mathcal{B}$  has subtensors, defined as setting the  $n$ th index to a value  $\alpha$ , that are orthogonal ([De Lathauwer et al., 2000](#)). Additionally, the subtensors have a descending order of Frobenius norms:

$$\|\mathcal{B}_{i_n=1}\| \geq \|\mathcal{B}_{i_n=2}\| \geq \cdots \geq \|\mathcal{B}_{i_n=I_n}\| \geq 0.$$

These Frobenius norms are the  $n$ -mode singular values of  $\mathcal{A}$ , and the vector  $\mathbf{U}_i(n)$  is the  $i$ th  $n$ -mode singular vector.

### Tensor-Train decomposition

The Tensor-Train decomposition factorizes an  $N$ -th order tensor  $\mathcal{A}$  in a product of third-order tensors ([Oseledets, 2011](#)). These are denoted as  $\mathbf{U}_1, \mathcal{U}_2, \dots, \mathbf{U}_N$  and are called TT-cores:

$$\mathcal{A}(i_1, \dots, i_N) = \mathbf{U}_1(i_1, :) \mathcal{U}_2(:, i_2, :) \mathcal{U}_3(:, i_3, :) \cdots \mathbf{U}_N(:, i_N).$$

The package computes this decomposition using the TT-SVD algorithm. The user can provide the ranks directly or allow them to be selected using an approximation accuracy.

## 5 Model fitting and inference

One major functionality that `tensoModels` provides is MLE estimation. This is done via two different algorithms: for the tensor variate normal distribution, a flip-flop algorithm is used. For tensor variate skewed distributions, an EM algorithm is used.

Before discussing the tensor variate normal, we discuss the simpler matrix variate normal distribution. The matrix variate normal distribution can be considered a special case of the multivariate normal distribution, where we force the restriction that the covariance matrix is separable. That is, the covariance of the vectorized matrix draws can be described as the Kronecker product of two matrices.

Let  $\mathbf{X} \sim N(\mathbf{M}, \mathbf{U}, \mathbf{V}) \in \mathbb{R}^{n \times p}$  and set  $\mathbf{y} = \text{vec}(\mathbf{X})$ . The matrix variate normal distribution imposes the separable covariance restriction

$$\mathbf{\Sigma} = \mathbf{V} \otimes \mathbf{U},$$

where  $\mathbf{U}$  is the row covariance matrix and  $\mathbf{V}$  is the column covariance matrix. Therefore

$$\text{vec}(\mathbf{X}) \sim N_{np}(\text{vec}(\mathbf{M}), \mathbf{V} \otimes \mathbf{U}).$$

This is a nonlinear restriction on the natural parameter space of the full multivariate normal family. For the matrix variate normal, the covariance matrix is restricted to be defined as the Kronecker product of two covariance matrices. This makes its parameter space a lower-dimensional curved subset of the full multivariate normal distribution (Roś et al., 2016).

By default, the initial values for the parameters are chosen in different ways. The location parameter  $\mathcal{M}$  is started as the median of the draws. This is because the mean of the draws is dependent on the skew and parameters. Each covariance matrix is initialized by centering by the median, matricizing along the corresponding dimension, and adding up the scatter matrix across all of the draws. This total is normalized to have a trace equal to its dimension to ensure the scale is not disproportionate. Initial values for the parameters are set to be conservative and specific to the distribution. For example, for the skewed  $t$ ,  $\nu = 20$ . Finally, the skew parameter  $\mathcal{A}$  is computed as a function of the expected value of the draws,  $\mathcal{M}$ , and the relevant parameters for that distribution.

The individual covariance matrices in a Kronecker product are not defined uniquely. Multiplying one covariance matrix by a constant and another by the inverse of that constant produces the same full covariance matrix. The package addresses this by normalizing the trace of each covariance matrix except the final one to equal its corresponding mode dimension. This preserves the overall Kronecker covariance while giving the individual mode covariance estimates a consistent scale.

For the matrix variate normal model, differentiating the likelihood gives the updates

$$\hat{\mathbf{V}} = \frac{1}{Np} \sum_{i=1}^N \mathbf{x}_{ic}^T \hat{\mathbf{U}}^{-1} \mathbf{x}_{ic}$$


and

$$\hat{\mathbf{U}} = \frac{1}{Nn} \sum_{i=1}^N \mathbf{x}_{ic} \hat{\mathbf{V}}^{-1} \mathbf{x}_{ic}^T.$$

Each estimate depends on the other, which leads to the iterative flip-flop algorithm. If  $N > \max(n, p)$ , then the flip-flop algorithm has only one solution and the MLEs are unique under the conditions described by Srivastava et al. (2009).

Convergence checks are done by comparing the change in likelihood relative to the previous iteration. At the  $t$ -th iteration, the equation used is

$$A = \frac{\loglik_t - \loglik_{t-1}}{\loglik_{t-1}}$$



**Figure 2:** From left to right, the input data is a multivariate normal, matrix variate normal with a separable covariance structure, matrix variate with nonseparable covariance matrices, and tensor variate normal with a separable covariance structure. When there is a separable covariance structure, the tensor variate distances should be comparable to the vector distances, leading to a small  $D$  value and a large  $p$ -value.

The algorithm ends when  $A < \epsilon$ , and by default  $\epsilon = 1e - 6$ .

Explain estimation, initialization, convergence checks, uncertainty quantification, and any constraints or identifiability conventions.

## 6 Diagnostic Tests

We assess tensor variate normality using the Mahalanobis distance. If  $\mathcal{X}$  follows an order- $O$  tensor variate normal distribution, then  $\text{vec}(\mathcal{X})$  follows a multivariate normal distribution with mean  $\text{vec}(\mathcal{M})$  and covariance matrix  $\Sigma_O \otimes \cdots \otimes \Sigma_1$ . This is why we can vectorize the tensor draws and compare the multivariate Mahalanobis distances to the tensor variate Mahalanobis distances. If we have properly reconstructed the structure using the MLEs  $\hat{\mathcal{M}}$  and  $\hat{\Sigma}_O, \cdots, \hat{\Sigma}_1$ , then the standardized draws should be from a standard normal.

If a tensor normal covariance structure is present, then the multivariate and tensor variate distances should be close to one another. Following the matrix variate normality paper, this can be visualized with a distance-distance plot: when the tensor normal structure is appropriate, the distances lie close to the reference line; when the Kronecker product structure is absent, the distances diverge.

The function `mahalanobis_dist()` computes the tensor Mahalanobis distances using fitted MLEs, while `mahalanobis_compare()` computes both the tensor and vectorized distances. The function `mahalanobis_test()` performs a two-sample Kolmogorov–Smirnov test to compare their empirical distributions.

The figure below shows the Mahalanobis distance-distance plots along with the KS test statistic and  $p$ -value.

We are also interested in performing a likelihood ratio test to see if one of the covariance matrices might be the identity matrix. Thus, we run the algorithm and restrict these matrices.

**Table 2:** LRT for Full Tensor Variate

| restrict | statistic | df | p-val |
|----------|-----------|----|-------|
| None     | 0.00      | 0  | 1     |
| 1        | 13257.02  | 2  | 0     |
| 2        | 61541.52  | 5  | 0     |
| 3        | 163633.38 | 9  | 0     |
| 1, 2     | 75108.59  | 7  | 0     |
| 1, 3     | 178430.32 | 11 | 0     |
| 2, 3     | 226430.11 | 14 | 0     |

**Table 3:** LRT for Second Sigma Identity

| restrict | statistic | df | p-val |
|----------|-----------|----|-------|
| None     | 0.000     | 0  | 1.000 |
| 1        | 12103.247 | 2  | 0.000 |
| 2        | 9.138     | 5  | 0.104 |
| 3        | 47443.558 | 9  | 0.000 |
| 1, 2     | 12119.269 | 7  | 0.000 |
| 1, 3     | 59436.610 | 11 | 0.000 |
| 2, 3     | 47461.651 | 14 | 0.000 |

For a restriction set  $\mathcal{R} \subset \{1, \dots, D\}$ , the restricted model fixes  $\Sigma_r = I_{n_r}$  for every  $r \in \mathcal{R}$ , while all other covariance matrices are estimated. The likelihood ratio test compares

$$H_0 : \Sigma_r = I_{n_r} \forall r \in \mathcal{R} \quad \text{vs.} \quad H_A : \Sigma_r \neq I_{n_r} \text{ for at least one } r \in \mathcal{R}.$$

The test statistic is

$$2\{\ell(\hat{\theta}_{\text{full}}) - \ell(\hat{\theta}_{\mathcal{R}})\} \sim \chi^2_{\nu_{\mathcal{R}}}, \quad \nu_{\mathcal{R}} = \sum_{r \in \mathcal{R}} \left\{ \frac{n_r(n_r + 1)}{2} - 1 \right\}$$

If the  $p$ -value is large, then we fail to reject the null hypothesis and conclude the restricted model fits the data as well as the unrestricted model. If the  $p$ -value is small, we reject the null hypothesis, which states the unrestricted model has a significantly better fit than the restricted model.

As expected, all of the restrictions lead to a model that is significantly different from the full model. All the covariance matrices are not the diagonal to best model the draws.

```
res <- map_dfr(list(NULL, c(1), c(2), c(3), c(1, 2), c(1, 3), c(2, 3)),
  draws = tvn_draws, test_restrict)
```

```
knitr::kable(
  res,
  digits = 3,
  escape = FALSE,
  booktabs = TRUE,
  caption = "LRT for Full Tensor Variate")
```

Now, we generate data where the second covariance matrix is diagonal. The model that restricts the 2nd covariance to be diagonal has a  $p - val > 0.05$ , meaning it does not have a significantly smaller likelihood than the full unrestricted model.

## 7 Performance and validation

### 7.1 Parameter recovery

To validate the distributions, we can compare the MLEs to the true values used to generate draws.

```
mean((tvn_mles$mu - mu_true)^2)
#> [1] 0.04388826

for(i in 1:3) {
  true_scaled <- sigmas_true[[i]] / sum(diag(sigmas_true[[i]]))
  est_scaled <- tvn_mles$sigmas[[i]] / sum(diag(tvn_mles$sigmas[[i]]))


  print(mean((est_scaled - true_scaled)^2))
}
#> [1] 1.348412e-06
#> [1] 2.30351e-06
#> [1] 7.609639e-06
```

## 8 Examples

### 8.1 GRABMyo

The Gesture Recognition and Biometrics ElectroMyogram (GRABMyo) data set is a public collection of surface electromyography recordings from 43 participants ([Jiang et al., 2024](#)). Each participant performed 16 active hand gestures in seven repeated trials. Participants are the independent draws in this analysis, and each processed draw is a  $16 \times 28 \times 16$  gesture  $\times$  channel  $\times$  movement-phase tensor.

The complete download and preprocessing pipeline is provided in the supplementary script `scripts/grabmyo_preprocess.R`. It reads the public WFDB records, converts device counts to millivolts, removes the four unused channels, computes log-RMS amplitude over 16 equal phase bins, averages the seven trials within participant and gesture, and writes the compact processed data set loaded below.



We fit a tensor-variate normal model to the 43 participant tensors. For comparison, each tensor is matricized along each of its three modes, and an ordinary matrix-variate normal model is fitted to each dataset.

```
grabmyo_fit_timed <- function(label, draws, model = "normal") {
  message(
    "Fitting ", label, " with dimensions ",
    paste(dim(draws[[1]]), collapse = " x ")
  )
  started <- Sys.time()
  fit <- tensor_mle(
    draws,
    model = model,
    max_iter = 1000,
    tol = 1e-6,
    quiet = FALSE
  )
  fit$elapsed_seconds <- as.numeric(difftime(
    Sys.time(), started, units = "secs"
  ))
  fit
}

#fit the tensor
grabmyo_fits <- list(
  tensor = grabmyo_fit_timed("tensor", grabmyo_tensors)
)

for (mode in 1:3) { #fit the matricizations
  matrix_draws <- lapply(
    grabmyo_tensors,
```

**Table 4:** Tensor-normal and matrix-normal fits to the GRABMyo data.

| model  | dimensions   | loglik    | k      | AIC        | BIC        | secs | delta_AIC | delta_BIC |
|--------|--------------|-----------|--------|------------|------------|------|-----------|-----------|
| tensor | 16 x 28 x 16 | 637062.9  | 7844   | -1258437.8 | -1244623.0 | 1.8  | 0         | 0         |
| mat2   | 28 x 256     | -97528.3  | 40469  | 275994.6   | 347268.6   | 6.8  | 1534432   | 1591892   |
| mat3   | 16 x 448     | 19358.9   | 107879 | 177040.1   | 367036.6   | 48.0 | 1435478   | 1611660   |
| mat1   | 16 x 448     | -424861.5 | 107879 | 1065481.0  | 1255477.5  | 17.5 | 2323919   | 2500100   |


```

    matricization,
    k = mode
  )
  label <- paste0("mat", mode)
  grabmyo_fits[[label]] <- grabmyo_fit_timed(label, matrix_draws)
  rm(matrix_draws)
  gc(FALSE)
}
#> Iteration 50: criterion = 8.220e-06

```

All three matrix models attain larger maximized likelihoods because the tensor covariance family is nested within each matricized covariance family. BIC selects the tensor model, which estimates only 7,844 parameters, compared with 40,469 for the next-smallest matricization.

The estimated mode-specific correlation matrices provide separate summaries of dependence among gestures, channels, and movement-phase bins.



```

grabmyo_non_normal_models <- c(
  "skewt", "vargamma", "invgauss", "genhyper"
)

grabmyo_distribution_fits <- c(
  normal = list(grabmyo_fits$tensor),
  purrr::map(
    purrr::set_names(grabmyo_non_normal_models),
    \(model) grabmyo_fit_timed(

```

**Table 5:** Tensor distribution fits to the 43 processed GRABMyo participant tensors.

| model                   | loglik   | k     | AIC      | BIC      | secs | delta_AIC | delta_BIC |
|-------------------------|----------|-------|----------|----------|------|-----------|-----------|
| Skew-t                  | 651376.6 | 15013 | -1272727 | -1246286 | 7.3  | 0.0       | 0.0       |
| Generalized hyperbolic  | 651375.5 | 15014 | -1272723 | -1246280 | 7.1  | 4.2       | 5.9       |
| Normal-inverse-Gaussian | 651362.5 | 15013 | -1272699 | -1246258 | 7.2  | 28.1      | 28.1      |
| Variance-gamma          | 651260.9 | 15013 | -1272496 | -1246055 | 5.3  | 231.4     | 231.4     |
| Normal                  | 637062.9 | 7844  | -1258438 | -1244623 | 1.8  | 14289.3   | 1663.3    |

```

    label = model,
    draws = grabmyo_tensors,
    model = model
  )
)
)

```

Both information criteria select the skew- $t$  distribution, with the generalized hyperbolic distribution providing the next-best fit. The large improvement over the normal model remains after penalizing the non-normal models for estimating an additional  $16 \times 28 \times 16$  skewness tensor.

## 9 Discussion

The package `tensormodels` provides support for tensor variate distributions, allowing for random number generation, density calculation, and MLE estimation.

## 10 Reproducibility

```

#> R version: 4.5.2
#> Platform: aarch64-apple-darwin20
#> rjtools version: 1.0.21

```

## 11 Acknowledgments

Add funding, contributor, and institutional acknowledgments here.

## References

- B. W. Bader and T. G. Kolda. Algorithm 862: MATLAB tensor classes for fast algorithm prototyping. *ACM Transactions on Mathematical Software*, 32(4):635–653, 2006. doi: 10.1145/1186785.1186794. [p6]
- B. W. Bader and T. G. Kolda. Efficient MATLAB computations with sparse and factored tensors. *SIAM Journal on Scientific Computing*, 30(1):205–231, 2007. doi: 10.1137/060676489. [p6]
- B. W. Bader, T. G. Kolda, et al. *Tensor Toolbox for MATLAB*, 2025. URL <https://www.tensortoolbox.org/>. Version 3.8. [p6]
- O. Barndorff-Nielsen, J. Kent, and M. Sørensen. Normal variance-mean mixtures and z distributions. *International Statistical Review*, 50(2):145–159, 1982. [p5]



- J. D. Carroll and J.-J. Chang. Analysis of individual differences in multidimensional scaling via an n-way generalization of eckart-young decomposition. *Psychometrika*, 35(3):283–319, 1970. doi: 10.1007/BF02310791. [p9]
- L. De Lathauwer, B. De Moor, and J. Vandewalle. A multilinear singular value decomposition. *SIAM Journal on Matrix Analysis and Applications*, 21:1253–1278, 2000. doi: 10.1137/S0895479896305696. [p9]
- R. A. Harshman. Foundations of the parafac procedure: Models and conditions for an explanatory multimodal factor analysis. 1970. [p9]
- F. L. Hitchcock. The expression of a tensor or a polyadic as a sum of products. *Journal of Mathematics and Physics*, 6(1-4):164–189, 1927. doi: 10.1002/sapm192761164. [p9]
- N. Jiang, A. Pradhan, and J. He. Gesture recognition and biometrics electromyogram (GRABMyo). *PhysioNet*, June 2024. doi: 10.13026/89dm-f662. URL <https://doi.org/10.13026/89dm-f662>. Version 1.1.0. [p13]
- J. Kossaifi, Y. Panagakis, A. Anandkumar, and M. Pantic. Tensorly: Tensor learning in Python. *Journal of Machine Learning Research*, 20(26):1–6, 2019. URL <https://jmlr.org/papers/v20/18-277.html>. [p6]
- J. Li, J. Bien, and M. T. Wells. *rTensor: Tools for Tensor Analysis and Decomposition*, 2024a. URL <https://cran.r-project.org/package=rTensor>. R package version 1.4.9. [p6]
- J. Li, J. Bien, M. T. Wells, H. Xiao, and Y. Chen. *ttTensor: Tensor Train Decomposition*, 2024b. URL <https://cran.r-project.org/package=ttTensor>. R package version 1.0.1. [p6]
- I. V. Oseledets. Tensor-train decomposition. *SIAM Journal on Scientific Computing*, 33(5): 2295–2317, 2011. doi: 10.1137/090752286. [p9]
- L. Periša. *TensorToolbox.jl: Tensor Decompositions in Julia*, 2023. URL <https://github.com/lanaperisa/TensorToolbox.jl>. GitHub repository. [p6]
- pyttb developers. *pyttb: Python Tensor Toolbox*, 2026. URL <https://pyttb.readthedocs.io/>. Documentation accessed 2026-06-30. [p6]
- B. Roś, F. Bijma, J. C. de Munck, and M. C. M. de Gunst. Existence and uniqueness of the maximum likelihood estimator for models with a kronecker product covariance structure. *Journal of Multivariate Analysis*, 143:345–361, 2016. doi: 10.1016/j.jmva.2015.05.019. [p5, 10]
- M. S. Srivastava, T. von Rosen, and D. von Rosen. Estimation and testing in general multivariate linear models with kronecker product covariance structure. *Sankhyā: The Indian Journal of Statistics, Series A*, 71(2):137–163, 2009. [p10]
- TensorOperations.jl developers. *TensorOperations.jl*, 2026. URL <https://quantumkithub.github.io/TensorOperations.jl/stable/>. Documentation accessed 2026-06-30. [p6]

Anupama Kannan  
 Baylor University  
 Department of Statistics  
 Waco, Texas  
 ORCID: 0000-0003-2789-6997  
[Anupama\\_Kannan1@baylor.edu](mailto:Anupama_Kannan1@baylor.edu)

David Kahle  
 Baylor University  
 Department of Statistics  
 Waco, Texas

ORCID: 0000-0002-9999-1558

[David\\_Kahle@baylor.edu](mailto:David_Kahle@baylor.edu)

*Michael Gallagher*

*Baylor University*

*Department of Statistics*

*Waco, Texas*

ORCID: 0000-0002-5487-8965

[Michael\\_Gallagher@baylor.edu](mailto:Michael_Gallagher@baylor.edu)